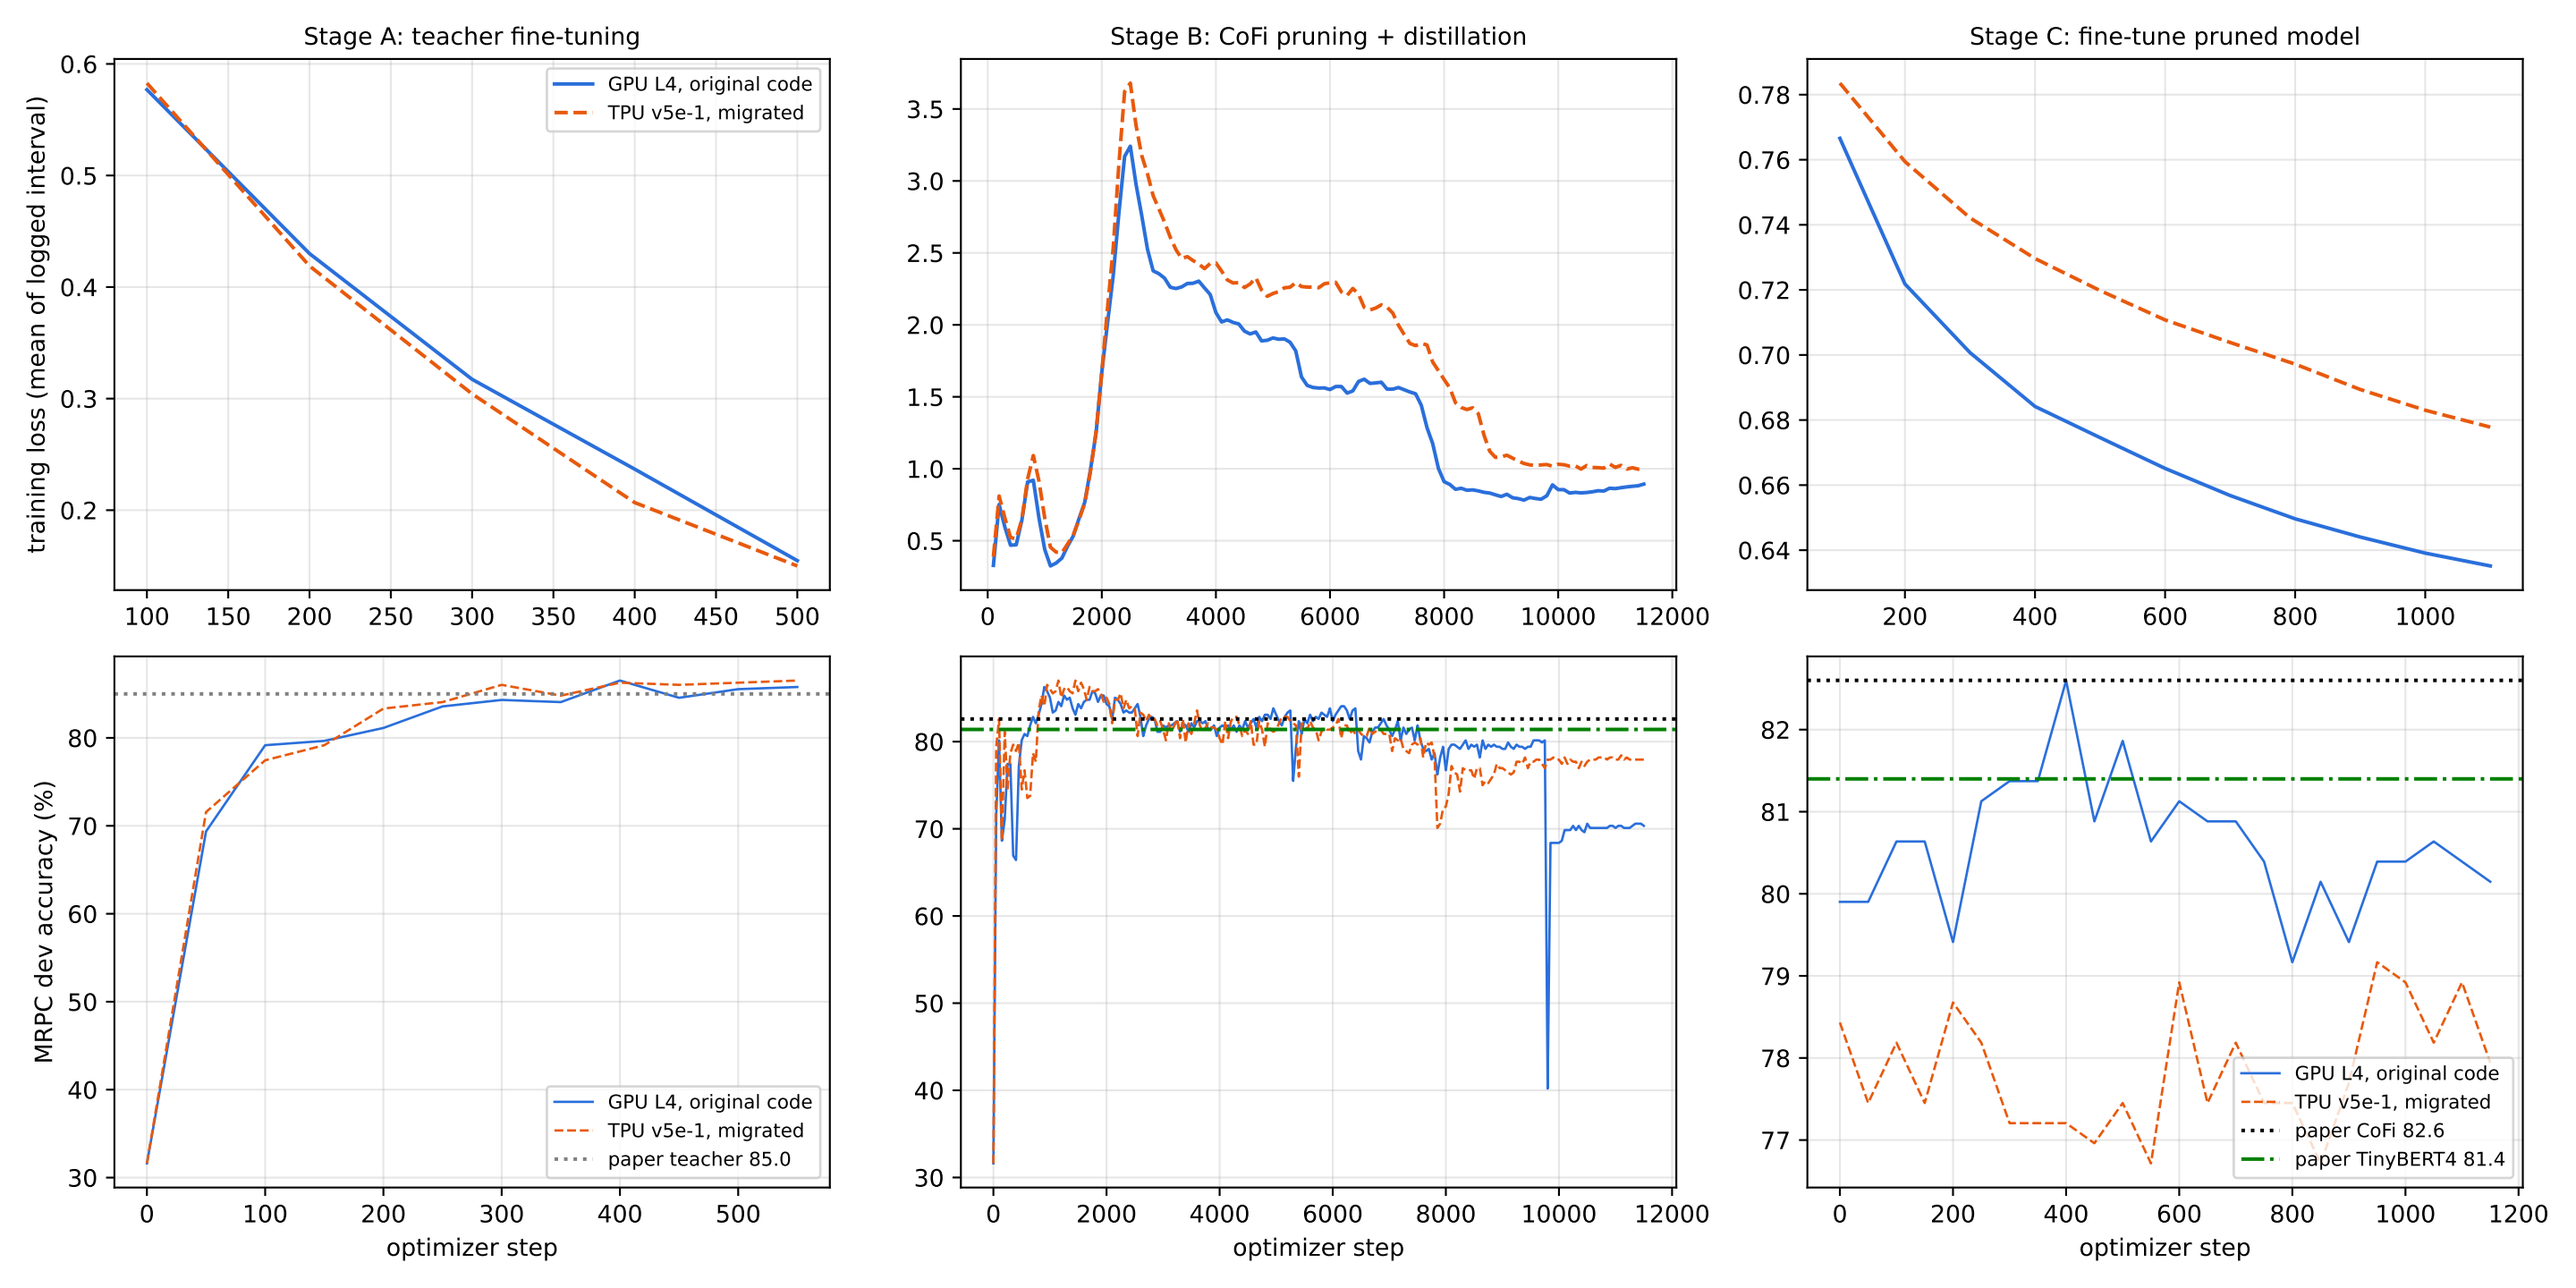
Stage A: teacher fine-tuning
Stage B: CoFi pruning + distillation
Stage C: fine-tune pruned model
training loss (mean of logged interval)
MRPC dev accuracy (%)
optimizer step
GPU L4, original code
TPU v5e-1, migrated
paper teacher 85.0
paper CoFi 82.6
paper TinyBERT4 81.4