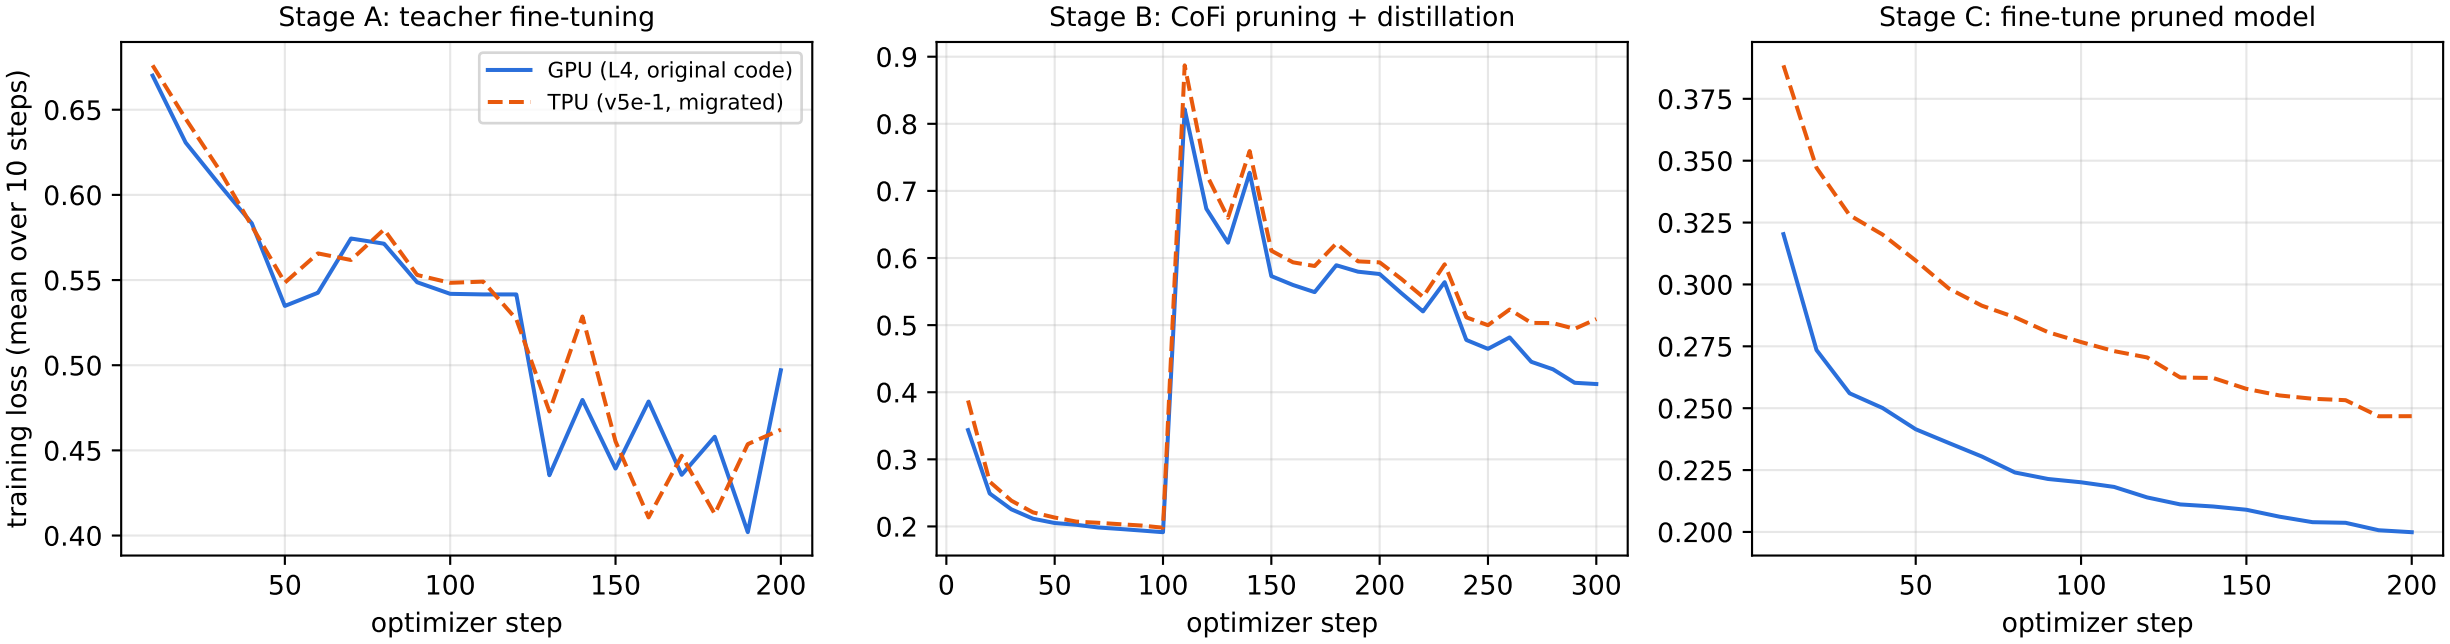
Stage A: teacher fine-tuning
Stage B: CoFi pruning + distillation
Stage C: fine-tune pruned model
GPU (L4, original code)
TPU (v5e-1, migrated)
training loss (mean over 10 steps)
optimizer step
optimizer step
optimizer step